



南京大學

NANJING UNIVERSITY

Introduction to

Algorithm Design and Analysis

[8] *log n* search



Yu Huang

<http://cs.nju.edu.cn/yuhuang>
Institute of Computer Software
Nanjing University



In the last class...

- **Selection – warm up**
 - Max and min
 - Second largest
- **Selection – rank k (median)**
 - Expected linear time
 - Worst-case linear time
- **Adversary argument**
 - Lower bound



The Searching Problem

- **Searching vs. Selection**
 - *Search* for “Alice” or “Bob”
 - The key itself matters
 - *Select* the “rank 2” student
 - The partial order relation matters
- **Expected cost for searching**
 - Brute force case: $O(n)$
 - Ideal case: $O(1)$
 - Can we achieve $O(\log n)$?



The Searching Problem


- Essential of *searching*
 - How to *organize the data* to enable efficient search
 - *logn* search
 - Each search cuts off half of the search space
 - How to organize the data to enable *logn* search
- *logn* search techniques
 - Warmup
 - Binary search over *sorted* sequences
 - *Balanced* Binary Search Tree (BST)
 - Red-black tree



Binary Search by Example

- **Binary search for “24”**
 - Divide the search space
 - Cut off half the space after each search

The sequence is
already sorted



Pseudo code in
p.129 [Baase01]

Binary Search Generalized

- **Peak-number**
 - Uni-modal array
- **Least number**
 - Not in the array
- **$A[i]=i$**
 - Sorted array




Balanced Binary Search Tree

- **Binary search tree (BST)**
 - Definitions and basic operations
- **Definition of Red-Black Tree (RBT)**
 - Black height
- **RBT operations**
 - Insertion into a red-black tree
 - Deletion from a red-black tree




Binary Search Tree Revisited

Good balancing
 $\Theta(\log n)$



Poor balancing
 $\Theta(n)$




 *In a properly drawn tree, pushing forward to get the ordered list.*

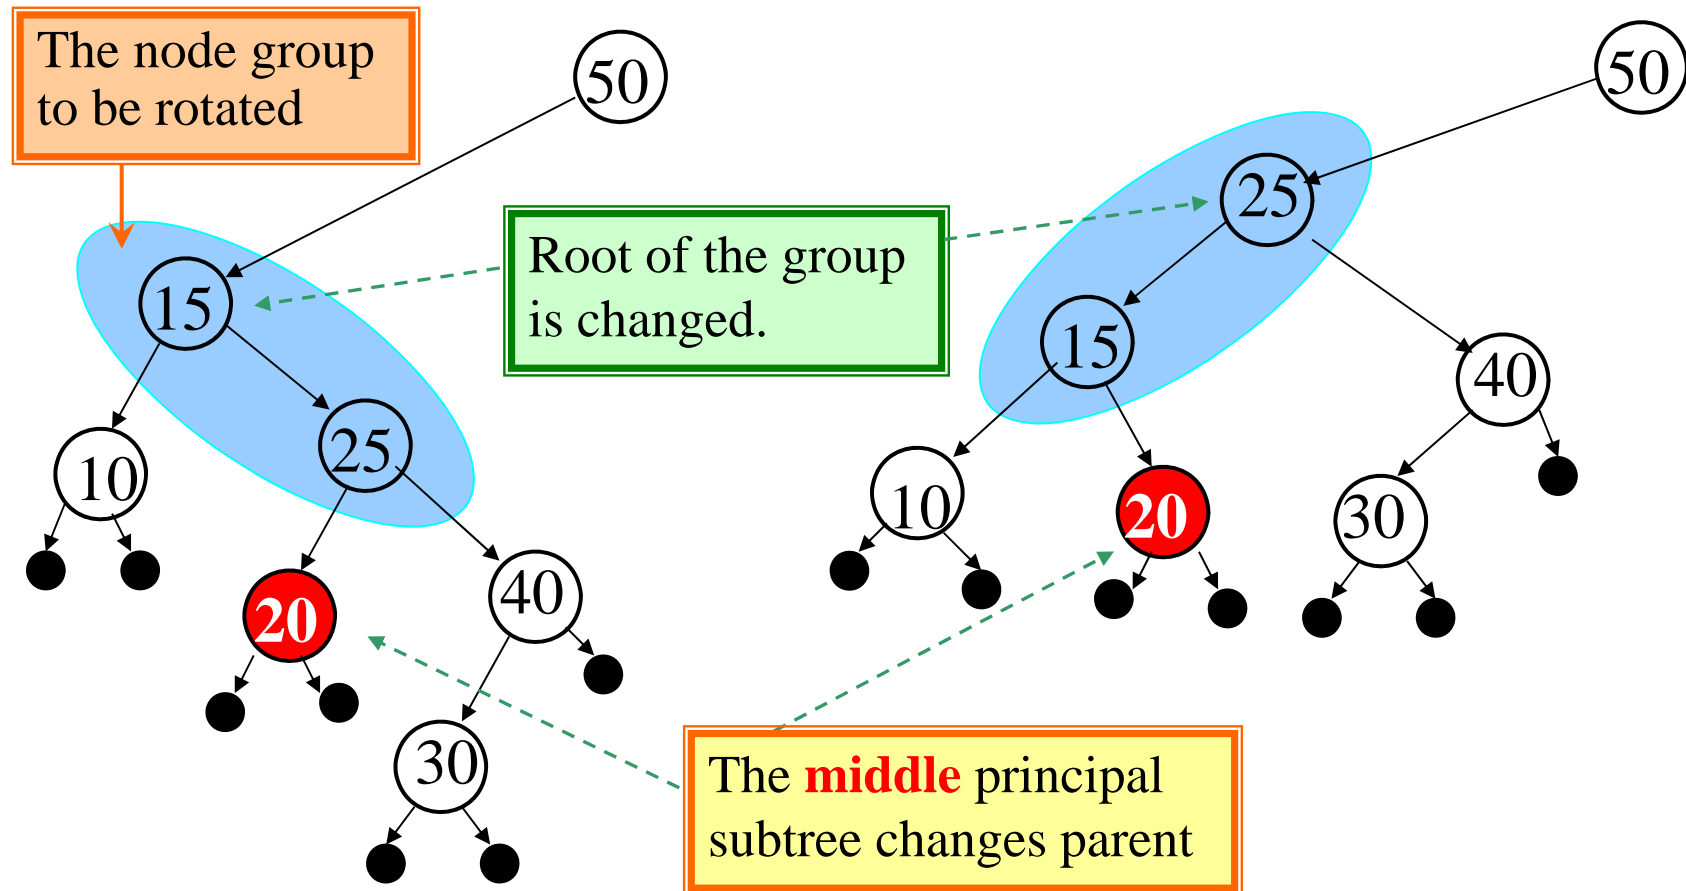
- Each node has a key, belonging to a linear ordered set
- An inorder traversal produces a sorted list of the keys

Node Group

As in 2-tree,
the number of
external node
is one more
than that of
internal node



Balancing by Rotation



Red-Black Tree: Definition

- If T is a **binary search tree** in which each node has a color, red or black, and all external nodes are black, then T is a **red-black tree** if and only if:
 - [*Color constraint*] No red node has a red child
 - [*Black height constraint*] The **black length** of all external paths from a given node u is the same (the black height of u)
 - The root is black.
- ***Almost*-red-black tree (ARB tree)**
 - Root is red, satisfying the other constraints.


Balancing is under control



Recursive Definition of RBT

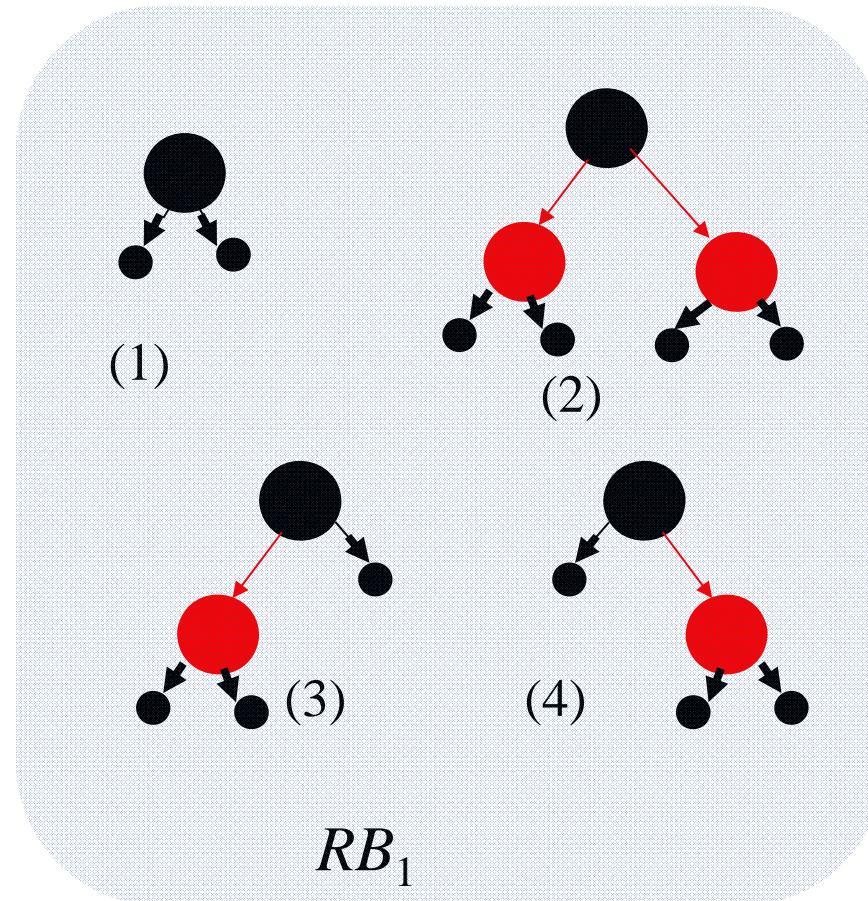


(A red-black tree of black height h is denoted as RB_h)

- **Definition:**



- An external node is an RB_0 tree, and the node is black.
- A binary tree is an ARB_h ($h \geq 1$) tree if: 
 - Its root is red, and
 - Its left and right subtrees are each an RB_{h-1} tree.
- A binary tree is an RB_h ($h \geq 1$) tree if:
 - Its root is black, and
 - Its left and right subtrees are each either an RB_{h-1} tree or an ARB_h tree.




RB_i and ARB_i




Red-Black Tree with 6 Nodes





Black edge



Black-depth Convention



All with the same
largest black depth: 2



ARB Trees



Properties of Red-Black Tree

- The **black height** of any RB_h tree or ARB_h tree is well-defined and is h .
- Let T be an RB_h tree, then:
 - T has at least $2^h - 1$ internal black nodes.
 - T has at most $4^h - 1$ internal nodes.
 - The depth of any black node is at most twice its black depth.
- Let A be an ARB_h tree, then:
 - A has at least $2^h - 2$ internal black nodes.
 - A has at most $(4^h)/2 - 1$ internal nodes.
 - The depth of any black node is at most twice its black depth.



Well-defined Black Height

- That “the **black height** of any RB_h tree or ARB_h tree is well defined” means *the black length of all external paths from the root is the same.*
- Proof: induction on h
- Base case: $h=0$, that is RB_0 (there is no ARB_0)
- In ARB_{h+1} , its two subtrees are both RB_h . Since the root is red, the black length of all external paths from the root is h , that's the same as its two subtrees.
- In RB_{h+1} :
 - Case 1: two subtrees are RB_h 's
 - Case 2: two subtrees are ARB_{h+1} 's
 - Case 3: one subtree is an RB_h (black height= h), and the another is an ARB_{h+1} (black height= $h+1$)



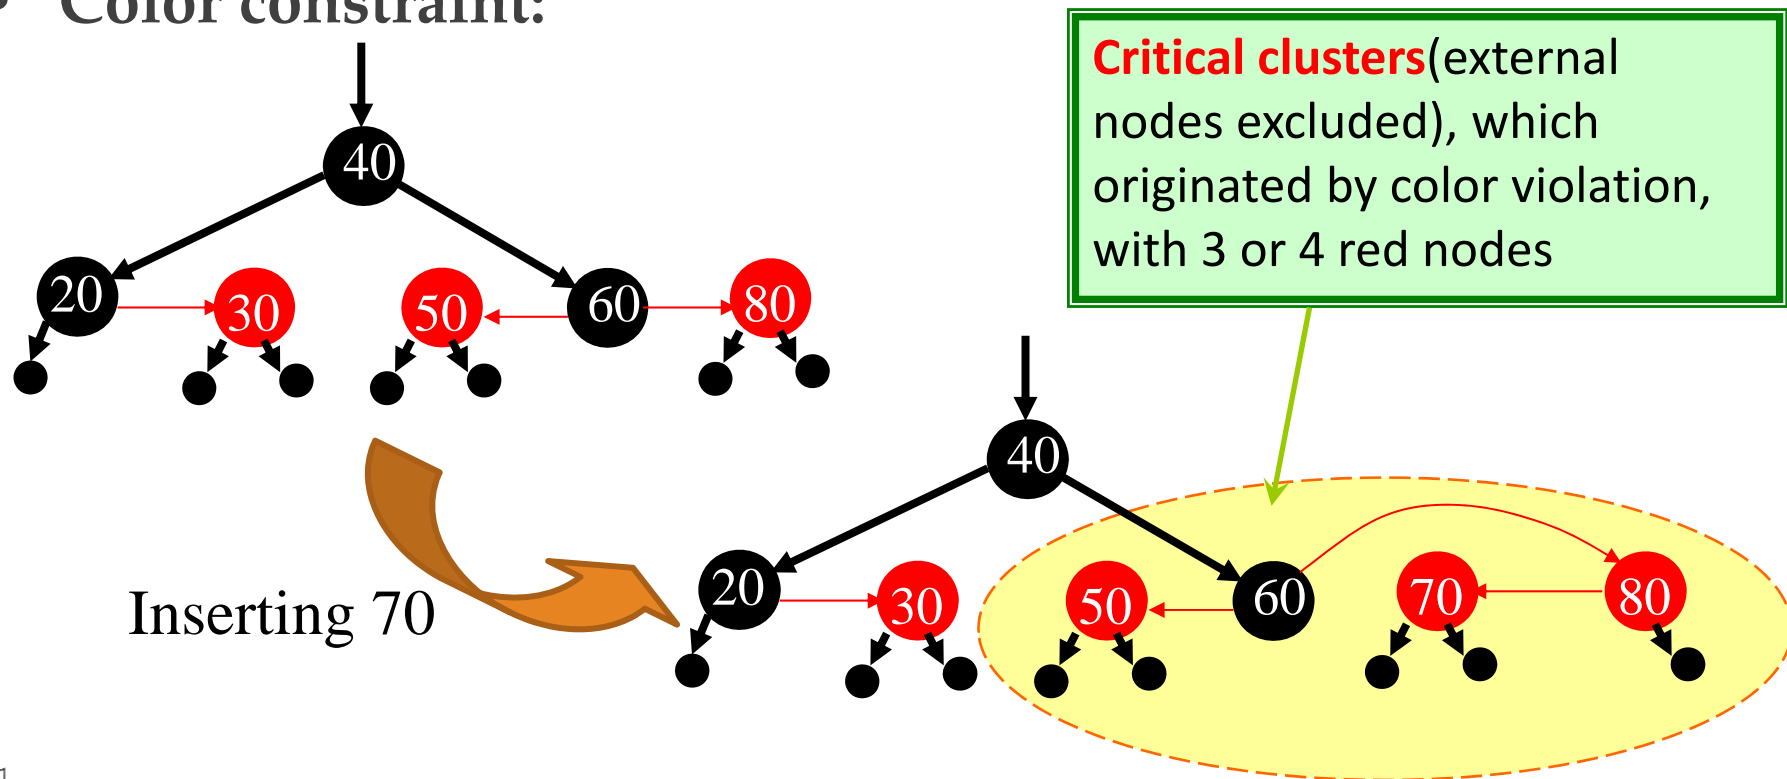
Bound on Depth of Node in RBTree

- Let T be a red-black tree with n internal nodes. Then no node has black depth greater than $\log(n+1)$, which means that the height of T in the usual sense is at most $2\log(n+1)$.
 - Proof:
 - Let h be the black height of T . The number of internal nodes, n , is at least the number of internal black nodes, which is at least $2^h - 1$, so $h \leq \log(n+1)$. The node with greatest depth is some external node. All external nodes are with black depth h . So, the depth is at most $2h$.




Influences of Insertion to an RBT


- Black height constraint:
 - No violation *if* inserting a red node.
- Color constraint:




Repairing 4-node Critical Cluster





Repairing 4-node Critical Cluster





New critical cluster with 3 nodes.
Color flip doesn't work,
Why?



Patterns of 3-node Critical Cluster




Shown as properly drawn




Repairing 3-Node Critical Cluster

Root of the critical cluster is changed to **M**, and the parentship is adjusted accordingly



The incurred critical cluster is of *pattern A*



Implementing Insertion: Class

```
class RBtree
```

```
    Element root;
```

```
    RBtree leftSubtree;
```

```
    RBtree rightSubtree;
```

```
    int color; /* red, black */
```

```
    static class InsReturn
```

```
        public RBtree newTree;
```

```
        public int status /* ok, rbr, brb, rrb, brr */
```

Color pattern



Implementing Insertion: Procedure

```
RBtree rbtInsert (RBtree oldRBtree, Element newNode)
```

```
  InsReturn ans; if (oldRBtree == nil) {
```

```
    ans.newTr = newNode;
```

```
    ans.newTr = newNode;
```

```
  } return ans.newTr;
```

the wrapper

```
  InsReturn rbtIns(RBtree oldRBtree, Element newNode)
```

```
  InsReturn ans, ansLeft, ansRight;
```

```
  if (oldRBtree == nil) then <Inserting simply>;
```

```
  else
```

```
    if (newNode.key < oldRBtree.root.key)
```

```
      ansLeft = rbtIns (oldRBtree.leftSubtree, newNode);
```

```
      ans = repairLeft(oldRBtree, ansLeft);
```

```
    else
```

```
      ansRight = rbtIns(oldRBtree.rightSubtree, newNode);
```

```
      ans = repairRight(oldRBtree, ansRight);
```

```
  return ans
```

the recursive function



Correctness of Insertion

- If the parameter `oldRBtree` of `rbtIns` is an RB_h tree or an ARB_{h+1} tree (which is true for the recursive calls on `rbtIns`), then the `newTree` and `status` fields returned are one of the following combinations:
 - `Status=ok`, and `newTree` is an RB_h or an ARB_{h+1} tree,
 - `Status=rbr`, and `newTree` is an RB_h ,
 - `Status=brb`, and `newTree` is an ARB_{h+1} tree,
 - `Status=rrb`, and `newTree.color=red`, `newTree.leftSubtree` is an ARB_{h+1} tree and `newTree.rightSubtree` is an RB_h tree,
 - `Status=brr`, and `newTree.color=red`, `newTree.rightSubtree` is an ARB_{h+1} tree and `newTree.leftSubtree` is an RB_h tree
- For those cases with red root, the color will be changed to black, with other constraints satisfied by repairing subroutines.





Deletion: Logical and Structural

u : to be deleted logically

π is parent of σ

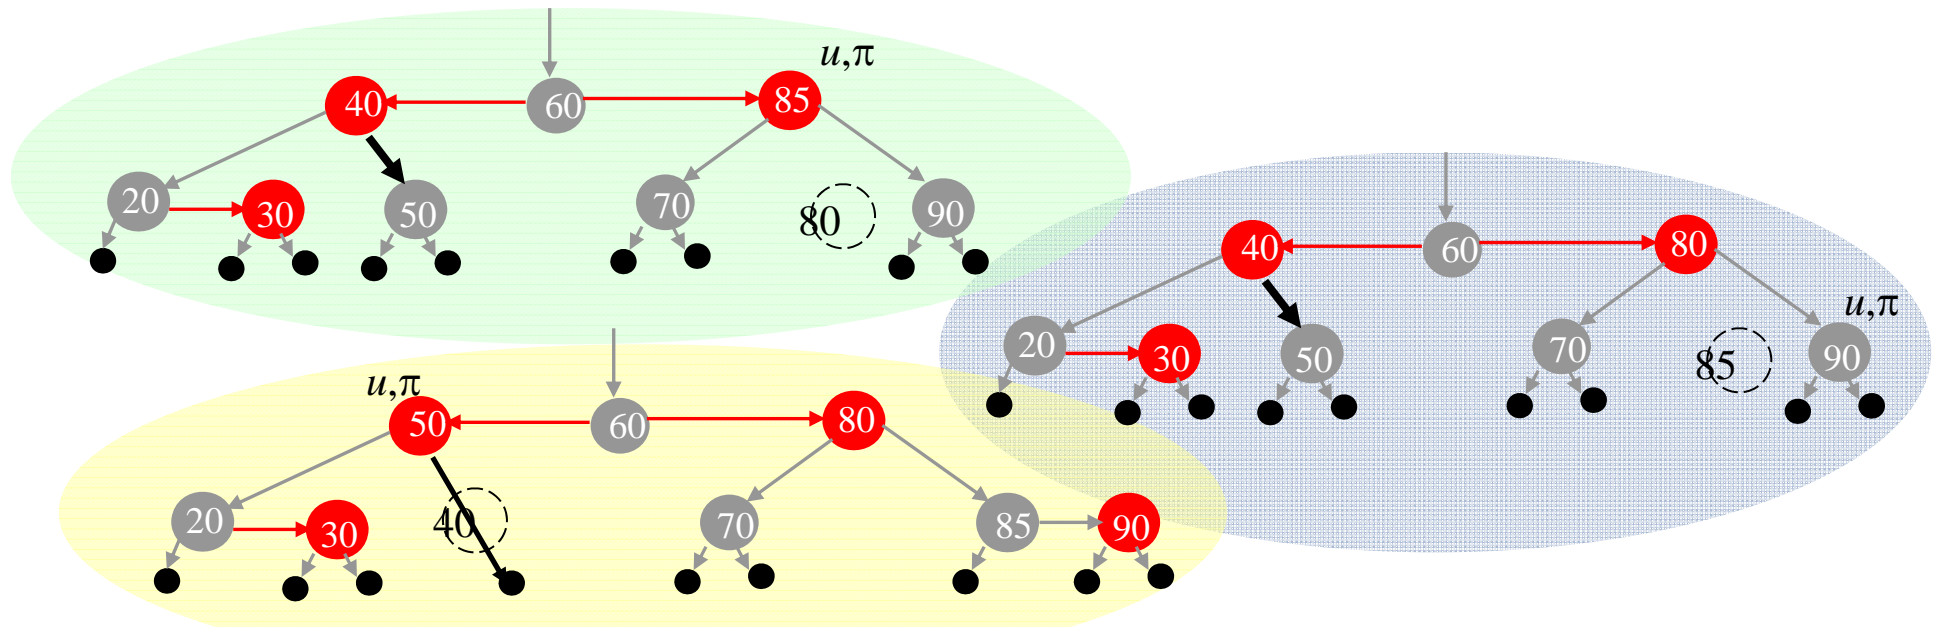

σ : tree successor of u , to be deleted structurally, with information moved into u

right subtree of S , to replace S

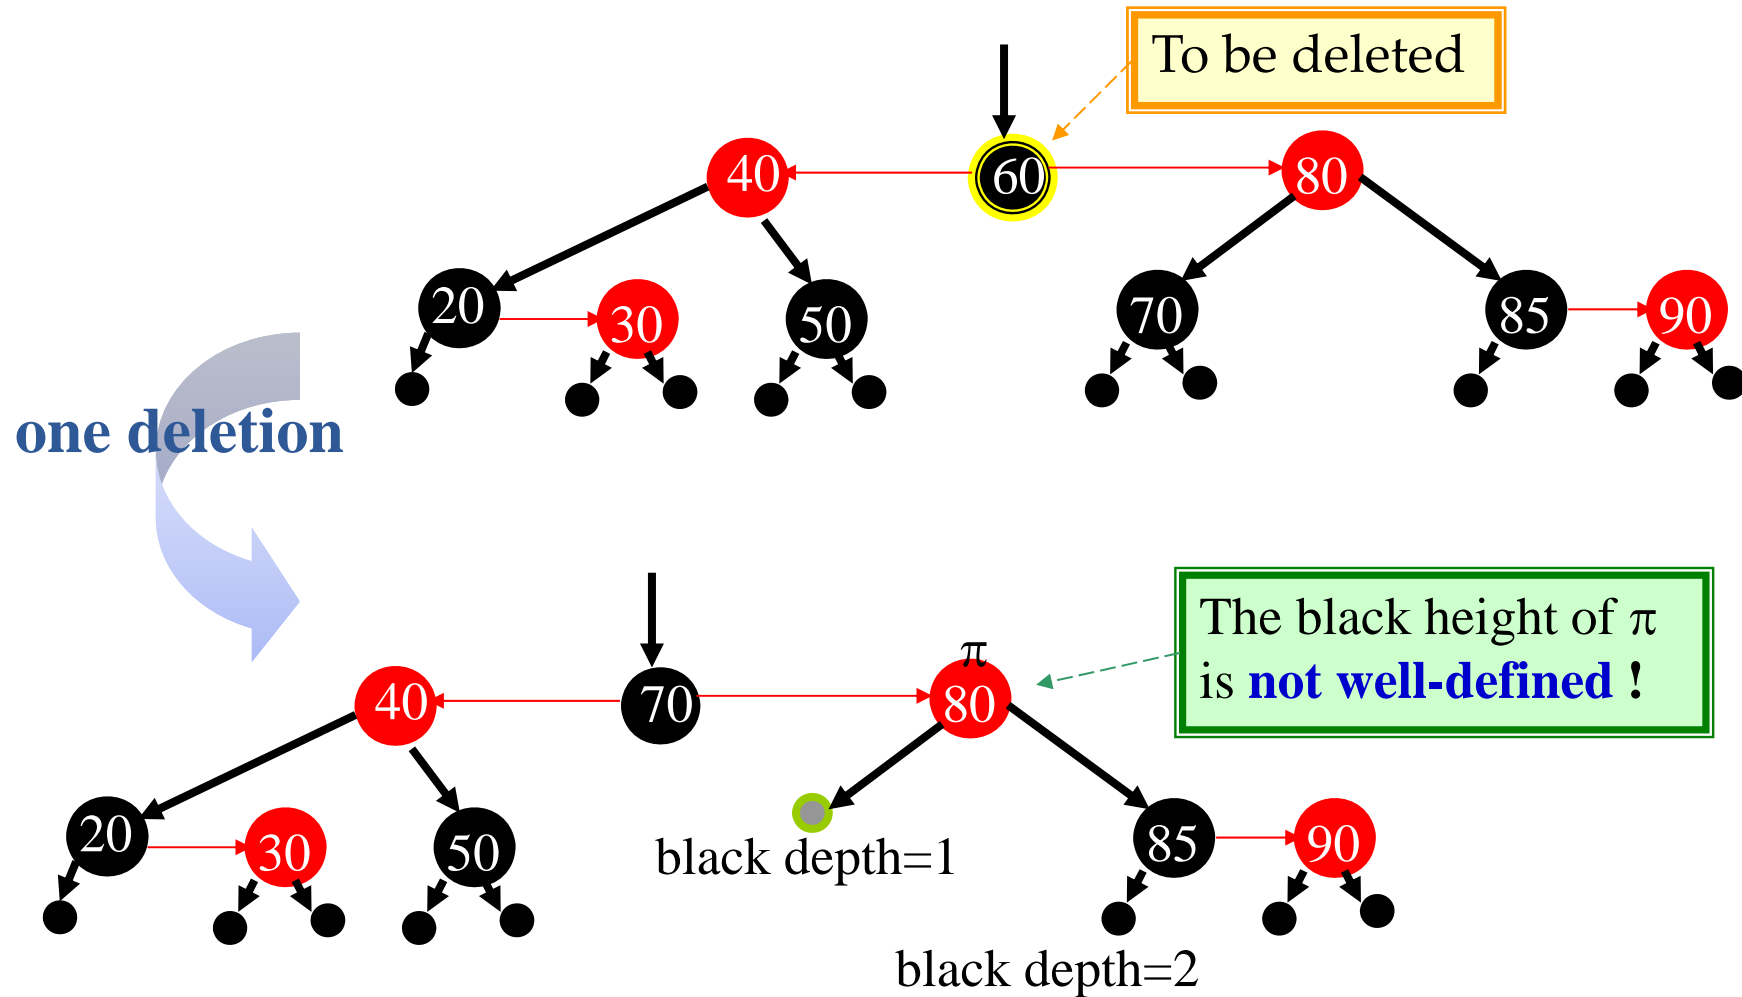


Deletion from RBT - Examples

Original tree



Deletion in RBT




Procedure of Red-Black Deletion


1. Do a standard BST search to locate the node to be logically deleted, call it u
2. If the right child of u is an external node, identify u as the node to be structurally deleted.
3. If the right child of u is an internal node, find the tree successor of u , call it σ , copy the key and information from σ to u . (color of u not changed) Identify σ as the node to be deleted structurally.
4. Carry out the structural deletion and repair any imbalance of black height.




Imbalance of Black Height



deleting 80




deleting 40




Black height has to be restored

deleting 85



deleting 60




Analysis of Black Imbalance

- **The imbalance occurs when:**
 - A black node is deleted structurally, and
 - Its right subtree is black (external)
- **The result is:**
 - An RB_{h-1} occupies the position of an RB_h as required by its parent, coloring it as a “gray” node.
- **Solution:**
 - Find a red node and turn it black as locally as possible.
 - The gray color might propagate up the tree.

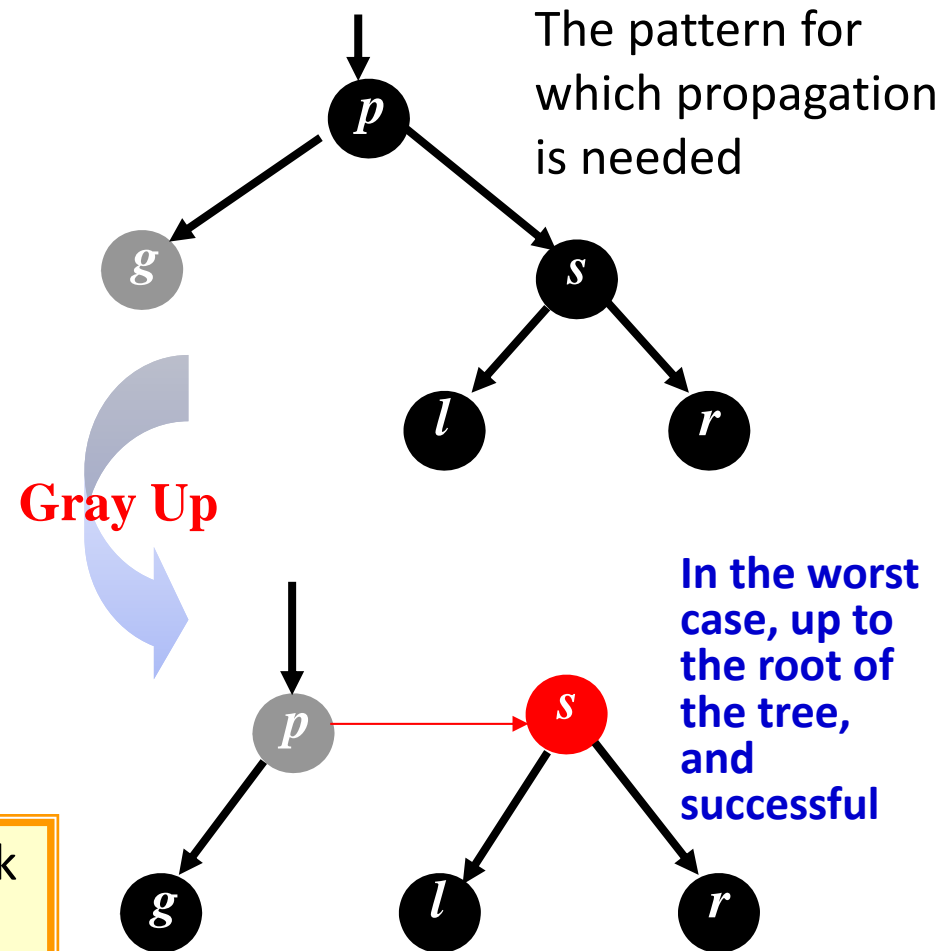


Propagation of Gray Node

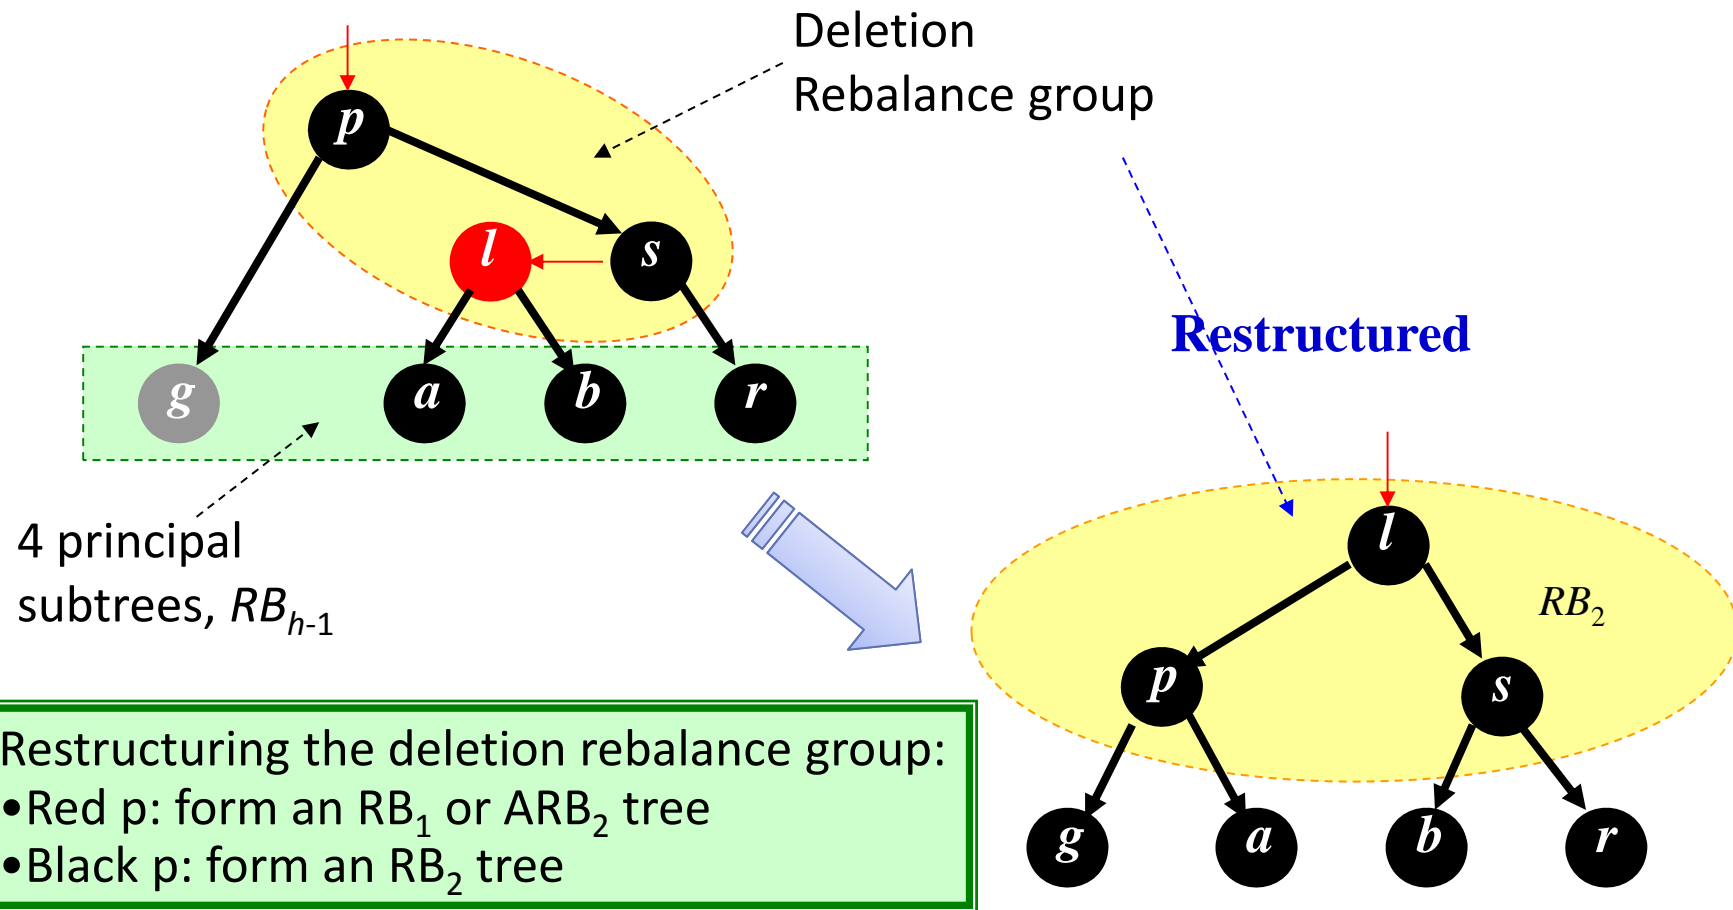


Map of the vicinity of g , the gray node

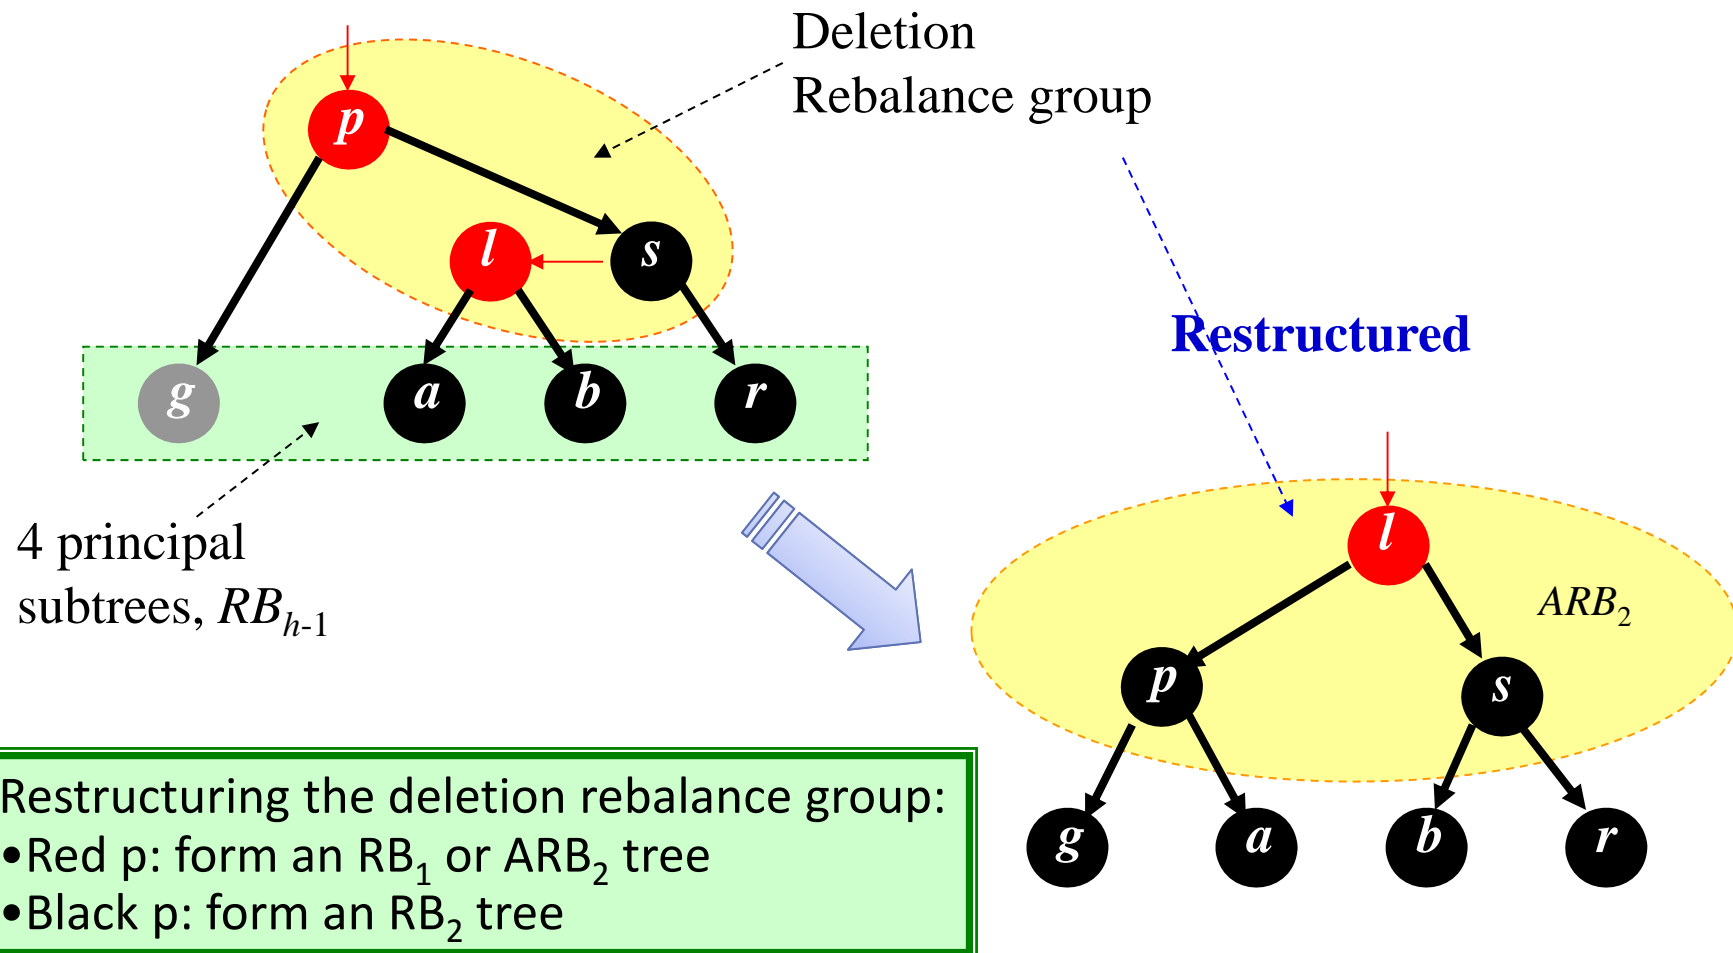
g -subtree gets well-defined black height, but that is less than that required by its parent



Repairing without Propagation



Repairing without Propagation



Complexity of Operations on RBT

- **With reasonable implementation**
 - A new node can be inserted correctly in a red-black tree with n nodes in $\Theta(\log n)$ time in the worst case.
 - Repairs for deletion do $O(1)$ structural changes, but may do $O(\log n)$ color changes.



Thank you!

Q & A

Yu Huang

<http://cs.nju.edu.cn/yuhuang>

